

HW2 Writeup: Decoding

Anjali Narayan-Chen and Johnny Chang

nrynchn2@illinois.edu, jychang3@illinois.edu

1 Introduction

For this homework, we adapt the approach of (Langlais et al., 2007) of a local greedy search algorithm. Using a predefined set of operations, the greedy search proceeds by modifying the current-best translation returned by the default stack decoder and searching over the set of modified translations in a hill-climbing approach. If a better-scoring translation is found, the modified translation is returned as the system’s prediction.

2 Algorithm

Given *source*, a sentence to translate, the core of the greedy search algorithm proceeds as follows:

Algorithm 1 Greedy Search

```
1: current  $\leftarrow$  seed(source)
2: loop
3:   s_current  $\leftarrow$  score(current)
4:   s  $\leftarrow$  s_current
5:   for all h  $\in$  neighborhood(current) do
6:     c  $\leftarrow$  score(h)
7:     if c > s then
8:       s  $\leftarrow$  c
9:       best  $\leftarrow$  h
10:  if s = s_current then return current
11:  else
12:    current  $\leftarrow$  best
```

The idea behind the greedy search is to modify parts of the translation one at a time, exploring the search space of neighboring hypotheses in a hill-climbing fashion. The hope is that such modifica-

tions to the sentence output by the original stack decoder will provide search directions that lead to improved hypotheses.

In our implementation, the function `seed` that seeds the search with an initial state is the provided stack decoder. The scoring function, `score`, is similarly defined as the scoring function used in the provided code, $\log p(\mathbf{f}, \mathbf{a}|\mathbf{e}) + \log p(\mathbf{e})$. We next define the `neighborhood` function, which takes a candidate translation as an argument and returns a set of neighboring hypotheses to consider.

2.1 Neighborhood Function

The neighborhood function is defined via a set of six operations that can transform a current translation, as defined by (Langlais et al., 2007).

Swap The *swap* operation swaps two adjacent target segments. This is designed to combat the baseline model’s strong bias toward monotonous translations.

Replace Given a specific source segment, the *replace* operator exchanges the translation for that segment with another found in the phrase translation table.

Bi-Replace The *bi-replace* operation works similarly to *replace*, allowing the translation of two adjacent source phrases to change simultaneously. The motivation for this is to modify the sentence enough to escape a possible local maximum in the search.

Split The *split* operation splits a given source phrase into two parts and re-translates the split

phrases according to translations found in the phrase translation table.

Merge The *merge* operation is the opposite of the *split* operation: it allows two adjacent source phrases to be merged. The merged phrase then receives a new translation according to the phrase translation table.

3 Results

References

Philippe Langlais, Alexandre Patry, and Fabrizio Gotti.
2007. A greedy decoder for phrase-based statistical
machine translation. *Proc. of TMI*.